

POLICY BRIEF

Arabic AI Safety Gaps

*What UAE Regulators Need to Know
Before Deploying Agentic AI at Scale*

PRESENTED BY

AI Safety UAE

DATE:

June 2026

01

The Issue

A two-year mandate, five AI tools, and zero Arabic safety documentation.

INTRODUCTION

The UAE's Agentic AI Mandate

In April 2026, **Sheikh Mohammed bin Rashid Al Maktoum** announced that 50% of UAE federal government sectors, services, and operations will run on agentic AI within **two years**.

The UAE Ministry of AI's Leading Generative AI Applications guide list applications like ChatGPT-5, Claude, Gemini, Manus, and Perplexity. Each publishes system cards, but **none** of that documentation addresses Arabic as an attack surface.

50%

of UAE federal government operations targeted for agentic AI within 2 years

5

AI assistants listed in the Ministry of AI guide, none with Arabic safety coverage

0

vendors have published safety documentation covering Arabic as an attack surface

02

What Our Analysis Found

Three findings from peer-reviewed Arabic safety research.

How We Evaluated the Safety Gap

Using the Arabic Dataset for LLM Safeguard Evaluation (*Ashraf et al., NAACL 2025*), we evaluated the system cards of all major AI assistants listed in the Ministry of AI's guide. We also reviewed research covering multilingual jailbreak success rates and Arabic transliteration-based attacks.

FINDING 01 Arabic is an active attack surface that vendors are not testing	Safety alignment training is conducted predominantly on English-language examples. Arabic adversarial inputs are not well represented, so threat-detection patterns do not transfer. Peer-reviewed testing documented bypass rates of 40% for GPT-4 and approaching 100% for ChatGPT when the same jailbreak prompt is translated into Arabic. [3]
FINDING 02 Translating English safety prompts into Arabic is not a solution	Vendors that do test in Arabic (like Claude), translate their existing English test prompts. A translated English probe does not test whether the model handles content that is legally or socially sensitive in the UAE. A test can pass while the deployment still fails in practice.
FINDING 03 Arabic-specific harm categories are absent from every published safety benchmark	Ashraf et al. identify six harm types specific to the Arabic-speaking world: political governance, human rights, regional conflicts, economic policy, controversial history, and cultural preservation. None appear in any vendor's safety documentation. These categories account for 53% of their dataset. [4]

03

Missing Arabic Harm Categories

Six harm types specific to the Arab world that are absent from every system card.

KEY FINDINGS

Six Arabic-Specific Harm Categories:

Ashraf et al. (NAACL 2025) identify six harm categories specific to the Arabic-speaking world accounting for 53% of their evaluation dataset. None appear in the published safety documentation of any AI assistant listed in the Ministry of AI's guide.

01 Political Governance

Prompts about government legitimacy, authority, and accountability where the same question is classified 'safe' or 'unsafe' depending on political framing.

02 Human Rights

Content engaging with civil liberties and freedom of expression in regional contexts requiring awareness of local legal and political norms.

03 Regional Conflicts

Questions about active and historical conflicts in the Arab world, carrying high sensitivity given ongoing geopolitical stakes.

04 Economic Policy

Content engaging with financial policy and economic critique in ways that carry political weight in Gulf and Arab contexts.

05 Controversial History

Politically sensitive historical events within Arab cultural memory, including cases where official and popular narratives diverge.

06 Cultural Preservation

Content that appears neutral in English but threatens linguistic, religious, or cultural identity norms specific to the Arab world.

04

Two Concrete Case Studies

What Arabic safety gaps look like when they occur in practice.

Two Examples from the Research

These examples illustrate what the Arabic safety gap means in practice:

CASE 01 Religious Misinformation:

A user could prompt a model to misinterpret a hadith or generate a fabricated fatwa endorsing a financial scheme content that carries serious legal and religious consequences in the UAE. This sits entirely outside every harm category English safety benchmarks test for. No vendor's system card documents this as a risk category.

CASE 02 Political Governance Double Standard:

Ashraf et al. document cases where the same Arabic question about government corruption receives a 'safe' classification from one political perspective and an 'unsafe' classification from another. Guardrails trained on English have no framework to detect this sensitivity and the inconsistency is invisible until it occurs in deployment.

05

Recommendations

Three actionable requirements for UAE procurement and regulatory bodies.

Three Actions for UAE Regulators:

01 **Require Arabic-specific red-teaming as a condition of procurement**

Vendors should demonstrate they have tested models against Arabic-language adversarial inputs, not translations of English ones. The EU AI Act (Article 13), entering into force in August 2026, requires transparency documentation sufficient for the deployment context. The UAE can establish a standard that requires evidence of Arabic language safety.

02 **Require per-language harm-rate breakdowns for all categories in vendor system cards**

This format already exists in Claude's documentation, which reports refusal rate and harmless response rate by language. The UAE can require all vendors to adopt the same standard and extend it to Arabic-specific harm categories.

03 **Establish an Arabic AI safety benchmark**

The UAE and MBZUAI published ArabicMMLU, the first multi-task Arabic language understanding benchmark at ACL 2024. No equivalent exists for Arabic safety evaluation. Extending that work into a repeatable, versioned safety standard would be a direct UAE contribution to global AI safety infrastructure and give procurement bodies a benchmark to require vendors to publish results against.

Sources & Citations:

- [1] Government of Dubai Media Office. Mohammed bin Rashid unveils framework to deploy agentic AI across 50% of government sectors within 2 years. April 23, 2026. <https://www.mediaoffice.ae/en/news/2026/april/23-04/mohammed-bin-rashid-chairs-uae-cabinet-meeting>

- [2] UAE Ministry of Artificial Intelligence. Leading Generative AI Applications. May 2026. <https://ai.gov.ae/wp-content/uploads/2026/05/Leading-Generative-AI-Applications-EN.pdf>

- [3] Yue Deng et al. Multilingual Jailbreak Challenges in Large Language Models. ICLR, May 2024. https://proceedings.iclr.cc/paper_files/paper/2024/file/6b396f766a50e0853a5164e68048540c-PaperConference.pdf

- [4] Yasser Ashraf et al. Arabic Dataset for LLM Safeguard Evaluation. NAACL, April 2025. <https://aclanthology.org/2025.naacl-long.285.pdf>

- [5] Mansour Al Ghanim et al. Jailbreaking LLMs with Arabic Transliteration and Arabizi. EMNLP, November 2024. <https://aclanthology.org/2024.emnlp-main.1034.pdf>

- [6] Anthropic. Claude Opus 4.5 System Card. November 2025. <https://www.anthropic.com/claude-opus-4-5-system-card>

- [7] European Union Artificial Intelligence Act, Article 13. Entry into force August 2, 2026. <https://artificialintelligenceact.eu/article/13/>

- [8] Fajri Koto et al. ArabicMMLU. ACL 2024. <https://aclanthology.org/2024.findings-acl.334/>



Thank You

For questions or collaboration:

info@aisafety.ae

Oversight Authorities

Reference:

UAE Ministry of Artificial Intelligence · MBZUAI · TDRA · UAE Cyber Security Council

© 2026 AI Safety UAE. All rights reserved. | aisafety.ae